RACHIT SRIVASTAVA

AI Engineer

Bengaluru, India

+91-8168433423 · rachit.srivastava1994@gmail.com · linkedin.com/in/rachitsrivas · github.com/rachit1994

SUMMARY

Senior AI engineer with 10+ years building high-scale products and recent production work across RAG, agentic workflows, LLMOps, streaming inference, and model cost control. Selected impact includes 50M+ user-scale systems, sub-200ms retrieval targets, 35% LLM cost reduction, 3x throughput improvement. Strengths include production RAG, agentic workflows, LLMOps, streaming inference, evaluation pipelines, and model cost control.

TECHNICAL SKILLS

AI / LLM Engineering: RAG, LangChain, LangGraph, OpenAI API, Anthropic Claude, pgvector, Pinecone, LLM evals, prompt engineering
Backend: Python, Node.js, TypeScript, REST, GraphQL, Redis, PostgreSQL
Cloud: AWS, Docker, Kubernetes, CI/CD, CloudWatch
Frontend: React, Next.js, WebSockets, performance optimization

EXPERIENCE

AI Engineer | 2026 – Present

Happening Today

- Shipping production-grade LLM agents for event discovery — planning, reasoning, and tool use in LangGraph (planner, tool-calling, critic) with latency and quality targets tied to product metrics.
- Built end-to-end RAG pipelines: ingestion, chunking, embeddings (OpenAI text-embedding-3-large), pgvector semantic search, and reranking — sub-200ms retrieval at scale for real-time discovery.
- Streaming inference over Anthropic Claude (SSE + Node.js) with backpressure, token budgets, and graceful fallback on context-window overflow for live user sessions.
- Prompt architecture for multi-step agentic workflows — chain-of-thought planning, few-shot retrieval from vector store, structured outputs, and pre-release LLM evals (groundedness, instruction-following).

AI/ML Engineer | 2024 – 2025

MIRA

- Delivered production RAG and agentic AI for a 50M+ user GenAI surface — retrieval, grounding, and real-time generation under high-traffic, enterprise-style reliability requirements.
- Built LLM evaluation infrastructure: factual accuracy, groundedness, and instruction-following benchmarks; ROUGE + human-preference scoring; CI/CD regression gates to catch prompt drift pre-deploy.
- Designed embeddings pipelines and semantic retrieval — chunking strategy, vector search tuning, and ranking for high-recall knowledge access in production agent workflows.
- Cut LLM inference cost 35% via prompt compression, semantic similarity caching (Redis), and model routing (GPT-4o / GPT-4o-mini) while holding eval baselines and production SLOs.

Technical Lead / Senior Software Engineer | 2022 – 2024

Delta Exchange

- Designed and built core trading engine features for a real-time derivatives platform serving 5M+ users — order book rendering, WebSocket feed management, and margin calculation APIs.
- Refactored legacy service architecture to reduce inter-service latency, achieving a 60% system performance improvement measured end-to-end.
- Built a React component library and shared state management layer, cutting UI feature development time significantly and boosting conversion by 25%.
- Led technical debt reduction initiative: migrated critical services from callback-heavy Node to async/await patterns, reducing bug rate by 40%.

Software Engineer | 2019 – 2022

BetterPlace

- Built scalable SaaS modules (workforce management, compliance workflows) using React, Python (Django/Flask), and Node.js.
- Designed a batched data ingestion pipeline that increased system throughput by 3x, enabling the platform to onboard enterprise clients with large worker datasets.
- Introduced frontend performance patterns (lazy loading, memoisation, pagination) that reduced UI render times by 50%.
- Migrated monolith features to independently deployable microservices, improving deployment frequency and reducing rollback risk.

Software Developer | 2015 – 2019

Liftoff Pvt Ltd & Daffodils Software

- Delivered full stack features across web applications in early-career roles.
- Built REST APIs.
- Integrated third-party services, and shipped React/Angular frontend modules.